

System Design Document: CareerOS

Version: 1.0.0

Status: Draft / Implementation Ready

Project Type: Agentic AI / Career Management System

1. Executive Summary

CareerOS is an autonomous, agentic AI system designed to function as a long-term career operating system. Unlike static chatbots or simple job aggregators, CareerOS utilizes an **Agentic Cognitive Architecture** to Think, Plan, Act, and Reflect. It persistently tracks user progress, adapts to market feedback (e.g., rejections, interviews), and optimizes a user's career trajectory through continuous reinforcement learning loops.

2. Problem Statement & Scope

2.1 Core Challenge

Early-career professionals and students face a multivariate optimization problem:

- **Information Asymmetry:** Users do not know which skills yield the highest ROI in the current market.
- **Execution Paralysis:** Users struggle to convert long-term goals into daily executable tasks.
- **Feedback Loops:** Traditional platforms do not "learn" from user rejection; they simply present the next job listing.

2.2 Solution Scope

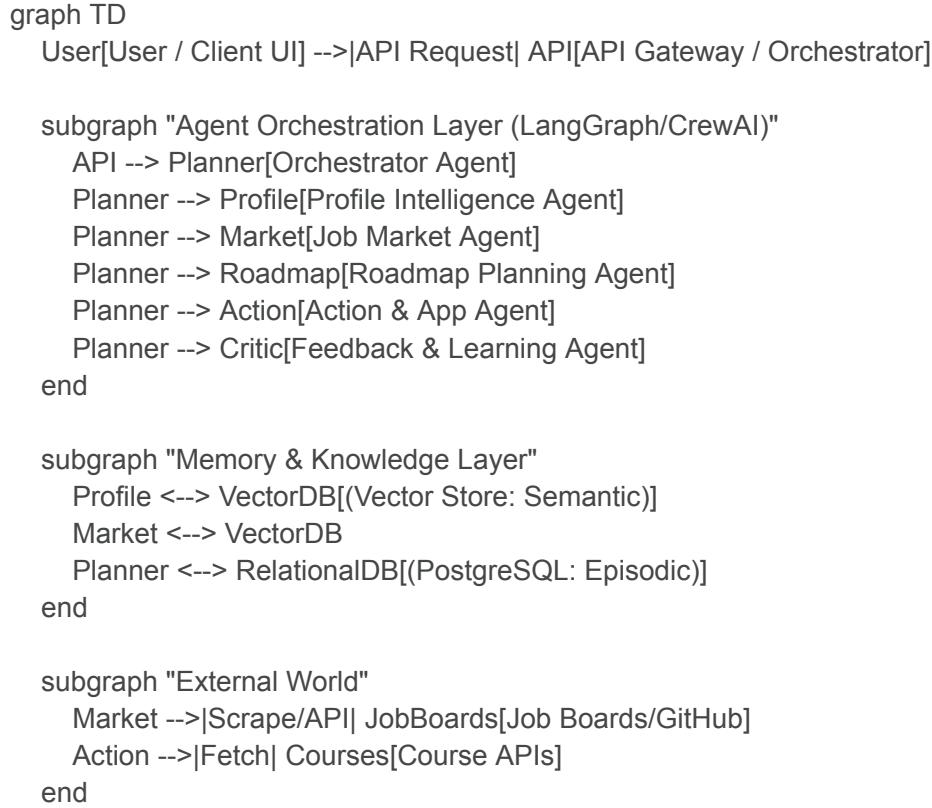
The system will:

1. Ingest and parse user artifacts (Resumes, GitHub, Portfolios).
2. Analyze real-time market data to identify skill gaps.
3. Orchestrate a multi-agent system to plan learning roadmaps and track applications.
4. Implement a memory architecture to store long-term context and adapt strategies based on outcomes.

3. High-Level System Architecture

The system follows a **Layered Agentic Architecture**, decoupling the user interface from the reasoning engines and memory stores.

3.1 Architectural Diagram



3.2 The Agentic Loop (Cognitive Cycle)

The system operates on a continuous control loop, formally defined as:

$\$\$State_{t+1} = f(State_t, Action_t, Observation_t)\$\$$

1. **Think:** Analyze current skills (S_t) vs. Market Demand.
2. **Plan:** Generate a Directed Acyclic Graph (DAG) of tasks (Learning Roadmap).
3. **Act:** Execute tools (Recommend Job, Draft Resume).
4. **Observe:** Monitor Application Status (Interview/Rejection).
5. **Reflect:** Update weights in the strategy (S_{t+1}) based on success/failure.

4. Component Design Specifications

4.1 Career Profile Intelligence Engine

- **Role:** The entry point for user data ingestion. Converts unstructured documents into structured JSON profiles.
- **Input:** PDF/DOCX Resume, GitHub URL, LinkedIn export.
- **Process:**
 - **Text Extraction:** PyMuPDF for layout-preserved extraction.
 - **NER (Named Entity Recognition):** spaCy model fine-tuned on skill taxonomies to extract specific technologies.

- **Embedding:** Sentence-Transformers to vectorize the user profile for semantic similarity matching.
- **Output JSON:**

```
{
  "profile_id": "uuid",
  "hard_skills": ["Python", "Docker", "FastAPI"],
  "soft_skills": ["Technical Writing", "Agile"],
  "experience_level": "Intermediate",
  "embedding_vector": [0.12, -0.45, 0.88, ...]
}
```

4.2 Job Market Reasoning Agent

- **Role:** Assesses the "value" of a skill or role in the current market.
- **Logic:**
 - Scrapes/ingests job descriptions (JDs) via Scrapy.
 - Calculates **Skill Gap Score**: The vector distance between `User_EMBEDDING` and `JD_EMBEDDING`.
 - **Demand Forecasting**: Analyzes frequency of keywords in JDs over the last 30 days.
- **Decision Gate:** If `Match_Score > 0.75` → Recommend Application. If `< 0.75` → Recommend Upskilling.

4.3 Personalized Skill Roadmap Planner

- **Role:** Deterministic planner that creates dependency graphs for learning.
- **Tools:** `LangGraph` for state management, `Neo4j` (optional) or `NetworkX` for dependency mapping.
- **Logic:**
 - Identifies missing skills (Set difference: `Target_{skills} - Current_{skills}`).
 - Queries "Learning Knowledge Base" for resources.
 - Output is a time-series schedule (e.g., "Week 1: Learn Pydantic").

4.4 Action Agent (Tracker & Application)

- **Role:** The "Hands" of the system.
- **Capabilities:**
 - **Resume Tailoring:** Uses LLM to rewrite bullet points to match JD keywords using `Jinja2` templates + `LLM-3`.
 - **Opportunity Finder:** Runs background cron jobs (`Celery`) to fetch new hackathons or internships matching the user's readiness score.

4.5 Continuous Learning & Feedback Agent (The Differentiator)

- **Role:** The "Meta-Learner." It does not execute tasks; it evaluates the performance of other agents.
- **Feedback Mechanism:**
 - **Trigger:** User logs an update (e.g., "Rejection from Google").
 - **Analysis:** The agent retrieves the resume used and the job description. It performs a "Post-Mortem" analysis to identify why the match failed (e.g., "Missing System Design experience").
 - **Action:** Updates the **Long-Term Memory** to de-prioritize similar roles until the skill gap is filled.

5. Data Architecture & Memory

To support "Agentic" behavior, we utilize a hybrid memory architecture.

5.1 Semantic Memory (Vector Database)

- **Technology:** ChromaDB or FAISS.
- **Content:**
 - Skill Embeddings (ESCo / O*NET taxonomy).
 - Job Description Embeddings.
 - Course/Project Content Embeddings.

5.2 Episodic Memory (Relational Database)

- **Technology:** PostgreSQL.
- **Content:**
 - **Application Logs:** History of every job applied to.
 - **Interactions:** User feedback, agent decisions.
 - **Outcomes:** Status (Applied, Interview, Rejected, Offer).

5.3 Procedural Memory (Rule Store)

- **Technology:** JSON/YAML Rules Engine.
 - **Content:** "If/Then" heuristics learned over time (e.g., "If user fails 3 Python interviews, trigger 'Advanced Python' module").
-

6. Technology Stack (Open Source / Free Tier)

Component	Technology Choice	Justification
LLM Inference	Ollama (LLaMA-3-8B / Mistral)	Local, free, high privacy.
Orchestration	LangGraph / CrewAI	State-of-the-art for cyclic agent flows.
Backend API	FastAPI (Python 3.11)	High performance, native async support.
Vector Store	ChromaDB	Open source, runs locally or in Docker.
NLP & parsing	spaCy + PyMuPDF	Industrial strength text processing.
Frontend	Next.js + Tailwind	React server components for fast UI.
Async Tasks	Celery + Redis	Background processing for scraping.
Deployment	Docker Compose	Containerization for easy reproducibility.

7. Security & Compliance

- **Data Minimization:** No scraping of login-walled platforms (LinkedIn) to avoid bans. Only public aggregation.
- **Encryption:** Resumes and personal data encrypted at rest (AES-256).

- **Local-First AI:** By using local LLMs (via Ollama), user career data does not leave the infrastructure, ensuring privacy.
-

8. Implementation Roadmap (Hackathon Strategy)

1. **Phase 1: Foundation (Hours 0-12)**
 - Set up FastAPI + ChromaDB.
 - Implement Resume Parsing & Profile Generation.
 2. **Phase 2: Intelligence (Hours 12-24)**
 - Implement Job Scraper & Matching Logic.
 - Build the Roadmap Generator.
 3. **Phase 3: The "Agent" Loop (Hours 24-36)**
 - Connect the Feedback Loop (Rejection \rightarrow New Plan).
 - Build the Frontend Dashboard.
 4. **Phase 4: Polish (Hours 36-48)**
 - Resume Tailoring output.
 - Final testing and "Why it wins" demo script.
-